

```
In [ ]: import pandas as pd
import numpy as np
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
import re
from gensim.models import KeyedVectors
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
stop_words = stopwords.words('english')
```

[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\erkhe\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!

```
In [ ]: def parse_data(file_path):
    with open(file_path, 'r', encoding='utf-8') as file:
        lines = file.readlines()

    data = []
    current_id = ''
    for line in lines:
        if line.startswith('###'):
            current_id = line.strip().rstrip('###')
        elif line.strip():
            parts = line.strip().split('\t')
            if len(parts) == 2:
                category, text = parts
                data.append({'id': current_id, 'category': category, 'text': text})

    return pd.DataFrame(data)

def clean_text(text):
    text = re.sub(r'[^a-zA-Z\s]', '', text, re.I|re.A)
    text = text.lower()
    text = text.strip()
    return text
```

```
In [ ]: df_train = parse_data('train.txt')
df_test = parse_data('test.txt')

df_train['text_cleaned'] = df_train['text'].apply(clean_text)
df_test['text_cleaned'] = df_test['text'].apply(clean_text)
```

```
In [ ]: vectorizer = TfidfVectorizer(stop_words=stop_words, max_features=3000, preprocessor=clean_text)
X_train = vectorizer.fit_transform(df_train['text_cleaned'])
y_train = df_train['category']
X_test = vectorizer.transform(df_test['text_cleaned'])
y_test = df_test['category']

lr_model = LogisticRegression(max_iter=1000, solver='sag', tol=0.1)
lr_model.fit(X_train, y_train)
lr_predictions = lr_model.predict(X_test)

print("Classification modèle bag of words + régression logistique :\n")
print(classification_report(y_test, lr_predictions))

cv_scores = cross_val_score(lr_model, X_train, y_train, cv=5, scoring='accuracy')
print("Validation croisée modèle bag of words, accuracy :", cv_scores.mean())
```

C:\Users\verkhe\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\sklearn\feature_extraction\text.py:409: UserWarning: Your stop_words may be inconsistent with your preprocessing. Tokenizing the stop words generated tokens ['arent', 'couldnt', 'didnt', 'doesnt', 'dont', 'hadnt', 'hasnt', 'havent', 'isnt', 'mightnt', 'mustnt', 'neednt', 'shant', 'shes', 'shouldnt', 'shouldve', 'thatll', 'wasnt', 'werent', 'wont', 'wouldnt', 'youd', 'youll', 'youre', 'youve'] not in stop_words.

Classification modèle bag of words + régression logistique :

	precision	recall	f1-score	support
BACKGROUND	0.58	0.49	0.53	2663
CONCLUSIONS	0.69	0.66	0.67	4426
METHODS	0.81	0.89	0.85	9751
OBJECTIVE	0.70	0.58	0.63	2377
RESULTS	0.82	0.83	0.83	10276
accuracy			0.77	29493
macro avg	0.72	0.69	0.70	29493
weighted avg	0.77	0.77	0.77	29493

Validation croisée modèle bag of words, accuracy : 0.7700890791529936

```
In [ ]: model_path = 'BioWordVec_PubMed_MIMICIII_d200.vec.bin'
embedding_model = KeyedVectors.load_word2vec_format(model_path, binary=True)

def sentence_vector(sentence, model):
    words = sentence.split()
    word_vectors = [model[word] for word in words if word in model]
    if len(word_vectors) == 0:
        return np.zeros(model.vector_size)
    else:
        return np.mean(word_vectors, axis=0)

X_train_embedded = np.array([sentence_vector(text, embedding_model) for text in df_train['text_cleaned']])
X_test_embedded = np.array([sentence_vector(text, embedding_model) for text in df_test['text_cleaned']])

lr_model_embedded = LogisticRegression(max_iter=1000)
lr_model_embedded.fit(X_train_embedded, y_train)
embedded_predictions = lr_model_embedded.predict(X_test_embedded)

print("Classification pour le modèle avec embeddings pré-entraînés :\n")
print(classification_report(y_test, embedded_predictions))
```

Classification pour le modèle avec embeddings pré-entraînés :

	precision	recall	f1-score	support
BACKGROUND	0.60	0.48	0.53	2663
CONCLUSIONS	0.68	0.70	0.69	4426
METHODS	0.82	0.87	0.85	9751
OBJECTIVE	0.65	0.61	0.63	2377
RESULTS	0.84	0.84	0.84	10276
accuracy			0.78	29493
macro avg	0.72	0.70	0.71	29493

weighted avg	0.77	0.78	0.77	29493
--------------	------	------	------	-------

In []: